

Text File IO

Class 17

Data

- variables are storage locations for data in RAM
- RAM is volatile
- its contents vanish when the program ends
- to make data persist across different runs of a program
- and across different programs
- we store data in files on disk

Files

- a file on disk is strictly a sequence of bytes
- when you ask the operating system for some stuff from a file, you just get raw bytes
- it is up to you, the programmer, how to **interpret** those bytes

Files

- a file on disk is strictly a sequence of bytes
- when you ask the operating system for some stuff from a file, you just get raw bytes
- it is up to you, the programmer, how to **interpret** those bytes
- there are two main flavors of file
 1. binary files
 2. text files
- **every** file is a binary file in the sense that it contains bytes
- text files, however, contain **only** bytes that correspond to ASCII characters
- one of those bytes represents the **newline** character, interpreted as the end of a line
- thus text files are easy to interpret as a sequence of **lines** each of which is a sequence of **characters**


Values

- a digital computer can only manipulate bits, 0s and 1s
- a group of 8 bits is a byte
- a byte can be interpreted as an unsigned integer value in the range 0 – 255
- thus, the only fundamental values a computer can manipulate are the 256 values in the range 0 – 255
- at the lowest level, there are no characters, or doubles, or negative numbers, or strings
- there are only the binary values 0 – 255

Characters

- to **represent** a character, a program must use an **encoding**
- an encoding is an agreement about which bit pattern will represent which character
- e.g., let us agree that in a character context the byte 0100 0001 (which is 65_{10} or $0x41$) will represent 'A'
- by default, C++ uses the ASCII encoding scheme, which we have seen several times

Text Files



what we see in an editor

A	l	i	c	e	\n	B	o	b	\n	C	a	r	o	l	\n
---	---	---	---	---	----	---	---	---	----	---	---	---	---	---	----

what is really in the disk file

- this is why **every** line of output must be terminated with a newline

Binary Files

- all files contain bytes
- the bytes encode some information, that we call data
- binary files contain data that is strictly designed to be read by computer programs
- some are open standard formats, e.g., jpeg and pdf for images
- some are proprietary, e.g., xls for spreadsheets and dwg for CAD drawings
- this is **not** the same as structured text files that conform to a format standard such as csv or xml

Text Files

- text files contain data that **can** be read either by a human or by a computer program
- working with text files requires the program to be able to
 - read** copy the data from the disk file **into** a program's variables
 - write** copy data from a program's variables **out** to some space on disk

Streams

- the model that C++ uses for working with files is to consider them as **streams** of bytes
- to **read** from a file is to treat the file as an **input** stream of bytes coming in from disk
- to **write** to a file is to treat the file as a destination for an **output** stream of bytes going out to disk

Streams

- the model that C++ uses for working with files is to consider them as **streams** of bytes
- to **read** from a file is to treat the file as an **input** stream of bytes coming in from disk
- to **write** to a file is to treat the file as a destination for an **output** stream of bytes going out to disk
- fortunately, we already know how to deal with streams
- for input, we use the stream extraction operator \gg and the function `getline`
- for output, we use the stream insertion operator \ll

Filenames

- files on disk are identified by **filename**
- by convention a filename consists of **name** and **extension**
- e.g., `grades.xls` or `phone_plan.cpp`
- by default, Windows file explorer does not show you the extension, because Windows was made for your grandma, not for programmers
- you should always turn on the listing of full file details so you can see file details, including the extension
- the extension indicates what **kind** of file it is
- there are hundreds of extensions
- we will use `.txt` for plain text files
- we will always assume that the text file is in the **current directory** (the project directory of CodeBlocks), so we will not have to worry about paths which are different on different operating systems

Opening an Input File

- to access a disk file it must first be **opened**
- to open a file means to associate it with a special **file variable**
- the file variable must be of type **ifstream** for an input text file
- this type is defined in the `<fstream>` library

```
#include <fstream>
...
ifstream input_file;
input_file.open("foo.bar");
```

- the identifier `input_file` is a programmer-defined name for **internal** use in the program
- the string `"foo.bar"` represents the **external** name of the file on disk

Opening an Output File

```
#include <fstream>
...
ofstream output_file;
output_file.open("foo.bar");
```

- the file variable must be of type **ofstream** for an output text file
- if the file **does not already exist**, it is created in the current directory

Opening an Output File

```
#include <fstream>
...
ofstream output_file;
output_file.open("foo.bar");
```

- the file variable must be of type **ofstream** for an output text file
- if the file **does not already exist**, it is created in the current directory
- **WARNING!** if the file **does** already exist, all its contents are **deleted** by the open function call

Closing a File

- before your program terminates
- you must **close** the file
- this frees up the operating system resources associated with the file
- for output files especially, this flushes the **write buffer** and ensures that everything you tried to write to the file is actually stored on disk

```
output_file.close();
```


Programs That Write to Files

- a newly opened file is empty
- each successive output operation appends more data after that already written
- cannot back up
- see files
 - Program 5-15 (page 274): write a few strings
 - Program 5-16: write strings without newlines (bad)
 - Program 5-17: numeric values entered from the keyboard, written to disk file
 - Program 5-18: string values entered from the keyboard, written to disk file

Programs That Read from Files

- the open function places the **read marker** at the first byte of the file
- as data are extracted from the file, the read marker is advanced toward the end of the file
- cannot back up

Programs That Read from Files

- the open function places the **read marker** at the first byte of the file
- as data are extracted from the file, the read marker is advanced toward the end of the file
- cannot back up
- if data are numeric, use stream extraction
`input_file >> value;`
- if data are strings, use getline to read an entire line
`getline (input_file, a_string);`
- see program 5-19, which uses getline instead of `>>` because it's reading strings

End of File

- a file may contain a little data
- or **a lot**
- typically we don't know how many lines of data a file contains
- the stream extraction operator `>>` **returns a value**
- it returns **true** if the read was successful
- it returns **false** if the read was not successful
- this allows us to control a while loop with the extraction operator
- see program 5-22

File Open Errors

- there are various conditions that may cause an **open** function to fail
 - attempt to read a file that does not exist
 - attempt to write (create) a file without OS permissions in that directory
 - attempt to write (create) a file on a disk that's full

File Open Errors

- there are various conditions that may cause an **open** function to fail
 - attempt to read a file that does not exist
 - attempt to write (create) a file without OS permissions in that directory
 - attempt to write (create) a file on a disk that's full
- in general, always need to **check** whether the open function succeeded
- do **not** use the method of program 5-23
- instead use method of middle of page 286 and program 5-24:

```
input_file.open(filename);  
if (!input_file.fail())  
{  
    ...  
}
```

Read and Write

- it is possible to open a file for **both** reading **and** writing at the same time
- but it is complex and uncommon
- for now, we will only allow reading **or** writing on any given file **but not both**

Read and Write

- it is possible to open a file for **both** reading **and** writing at the same time
- but it is complex and uncommon
- for now, we will only allow reading **or** writing on any given file **but not both**
- but it is completely acceptable to have **two** files open at the same time
 - one for reading
 - one for writing

A More Realistic Example

- the examples shown so far have been relatively simplistic
- as a more realistic example, consider a file structured as shown below
- for use as input to a program to generate phone plan invoices

A Rachel Carson

12.3

C Aldo Leopold

A Paul R. Erlich

7.25

B Farley Mowat

1.23

Omitted

- we are using C++11
- therefore the top half of page 288 is irrelevant
- expunge those nasty old C-strings whenever you can!

getline and Stream Extraction

- you are familiar with stream extraction \gg that skips whitespace and **stops at** either
 - whitespace
 - a character not valid for the variable being read (e.g., a decimal point for an integer)

getline and Stream Extraction

- you are familiar with stream extraction \gg that skips whitespace and **stops at** either
 - whitespace
 - a character not valid for the variable being read (e.g., a decimal point for an integer)
- you are familiar with getline that reads a **string** from the current position to the newline, skips the newline, and positions the read marker **at the beginning** of the next line

getline and Stream Extraction

- you are familiar with stream extraction \gg that skips whitespace and **stops at** either
 - whitespace
 - a character not valid for the variable being read (e.g., a decimal point for an integer)
- you are familiar with getline that reads a **string** from the current position to the newline, skips the newline, and positions the read marker **at the beginning** of the next line
- there is no problem with a getline **followed by** a stream extraction
- but stream extraction followed by a getline can cause **major problems**

Data Files

- consider a data file that has alternating lines of data
- the first line consists of name
- the second line has numbers

B	o	b	\n	
4	7	2	9	\n
A	n	n	e	\n
1	2	3	\n	

Getline and Stream Extraction

- we use `getline` for the first name
- no problem, the read marker reads B-o-b
- skips *newline*
- finishes on the 4

B	o	b	\n	
4	7	2	9	\n
A	n	n	e	\n
1	2	3	\n	

Getline and Stream Extraction

- we now use stream extraction \gg for the number
- no problem, we read 4-7-2-9 and the read marker ends on the *newline*

B	o	b	\n	
4	7	2	9	\n
A	n	n	e	\n
1	2	3	\n	

Getline and Stream Extraction

- now we want to use getline again to read Anne
- but the read marker is on the newline at the end of 4729
- getline reads until the next newline **but it's already on a newline**
- so it reads **0 characters**, skips the newline, and puts the read marker on the A of Anne

B	o	b	\n	
4	7	2	9	\n
A	n	n	e	\n
1	2	3	\n	

Getline and Stream Extraction

- now everything is hosed
- we're ready to read an integer
- but the read marker is on a character

B	o	b	\n	
4	7	2	9	\n
A	n	n	e	\n
1	2	3	\n	

Getline and Stream Extraction

- solution: after extraction \gg , we have to **skip over** the newline
- to read up to and skip over the newline, we use **getline** with a dummy variable: `getline(stream, dummy);`
- in programming a variable named **dummy** is used for a value that is never used
- a variable named `dummy` signals that its value will be ignored

```
getline(stream, string_variable);  
stream >> variable >> variable;  
getline(stream, dummy); // to consume the newline  
getline(stream, string_variable);  
stream >> variable >> variable;  
getline(stream, dummy); // to consume the newline  
...
```