

Searching

Class 27

Introduction

- two of the most fundamental concepts in computer science are, given an array of values:
 - **search** through the values to see if a specific value is present and, if so, where
 - **sort** the values in nondecreasing order

Introduction

- two of the most fundamental concepts in computer science are, given an array of values:
 - **search** through the values to see if a specific value is present and, if so, where
 - **sort** the values in nondecreasing order
- as a computer scientist, you must be able to understand and program several different algorithms for each of these tasks

Introduction

- two of the most fundamental concepts in computer science are, given an array of values:
 - **search** through the values to see if a specific value is present and, if so, where
 - **sort** the values in nondecreasing order
- as a computer scientist, you must be able to understand and program several different algorithms for each of these tasks
- in all of these slides, “array” is a generic term
- it means either an old-fashioned C-array or a modern STL vector

Searching

- an array that contains an arbitrary number of values
- each value is in a specific array location, but the values are not sorted

17	23	5	11	2	29	3
----	----	---	----	---	----	---

Searching

- an array that contains an arbitrary number of values
- each value is in a specific array location, but the values are not sorted

17	23	5	11	2	29	3
----	----	---	----	---	----	---

- search for the value 11
- do a **comparison** on 17, 23, 5, 11 (4 comparisons)
- return 3 (the position where 11 was found)

Searching

- an array that contains an arbitrary number of values
- each value is in a specific array location, but the values are not sorted

17	23	5	11	2	29	3
----	----	---	----	---	----	---

- search for the value 11
- do a **comparison** on 17, 23, 5, 11 (4 comparisons)
- return 3 (the position where 11 was found)
- search for the value 7
- do comparisons on 17, 23, 5, 11, 2, 29, 3 (7 comparisons)
- return a not-found indicator

Linear Search

- Program 8-1, page 464, a function to do this

```
int linearSearch(const int arr[], int size, int value);
```

- up to **size** elements of the array **arr** are searched for the existence of **value**
- if **value** is found within the first **size** elements of **arr**, the first position where value is encountered is returned by the function
- if value is not found, the sentinel value -1 is returned

Linear Search

- Program 8-1, page 464, a function to do this

```
int linearSearch(const int arr[], int size, int value);
```

- up to **size** elements of the array **arr** are searched for the existence of **value**
- if **value** is found within the first **size** elements of **arr**, the first position where value is encountered is returned by the function
- if value is not found, the sentinel value -1 is returned
- the search is called **linear search** because the algorithm searches along the **line** of array elements, one by one, starting at the beginning, trying to find value

Problems

- there are several problems with Gaddis' approach
- the biggest one is that none of his int data types will be accepted by a modern compiler like clang
- the data type for a size is `size_t`, and `size_t` is an **unsigned** integer; thus, it is **impossible** to return `-1` as a sentinel value for the not-found condition

Problems

- there are several problems with Gaddis' approach
- the biggest one is that none of his int data types will be accepted by a modern compiler like clang
- the data type for a size is `size_t`, and `size_t` is an **unsigned** integer; thus, it is **impossible** to return `-1` as a sentinel value for the not-found condition
- a solution to this problem is to return **the size of the array** as the not-found sentinel value
- in an array of size, say **10**, the largest valid index is **9**
- if the searched-for value is not found, we return 10, which is **not a valid index**, indicating not-found
- look at the complete program, using a vector:
`program_8_1_modified.cpp`

Notice

- several things to note in the modified program:
 - the return type of the function is `size_t`
 - the index and position are also of type `size_t`
 - there is no parameter for `size` because a vector knows how big it is

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector
 2. at random, pick a value in the range 0 – 10,000 and

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector
 2. at random, pick a value in the range 0 – 10,000 and
 - 2.1 see if it is in the vector or not

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector
 2. at random, pick a value in the range 0 – 10,000 and
 - 2.1 see if it is in the vector or not
 - 2.2 count how many comparisons it takes to find out

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector
 2. at random, pick a value in the range 0 – 10,000 and
 - 2.1 see if it is in the vector or not
 - 2.2 count how many comparisons it takes to find out
 3. repeat step 2 10,000 times

Linear Search Analysis

- it is traditional to use n to denote the number of elements of an array (this is the same thing as vector's size)
- do the following experiment (I encourage you to write this program):
 1. put 1,000 random values, chosen from 0 – 10,000, into a vector
 2. at random, pick a value in the range 0 – 10,000 and
 - 2.1 see if it is in the vector or not
 - 2.2 count how many comparisons it takes to find out
 3. repeat step 2 10,000 times
 4. calculate the average number of comparisons from step 2

Linear Search Analysis

- results:

Total hits: 930 Total misses: 9070

Minimum comparisons: 1 Maximum comparisons: 1000

Average number of comparisons: 954.747

- with 1,000 **random** values from 0 to 10,000
- the **average** number of comparisons is about 955
- have to search almost **the entire** array almost every time

Linear Search

Linear Search Pros

- very easy algorithm to understand
- easy algorithm to code correctly
- the only practical algorithm for **unsorted** values

Linear Search Cons

- inefficient for **sorted** values
- must examine $n/2$ elements on average for a value that **is** in the array
- must examine **all n elements** for a value **not** in the array

Linear Search

Linear Search Pros

- very easy algorithm to understand
- easy algorithm to code correctly
- the only practical algorithm for **unsorted** values
- for sorted values, we can do better

Linear Search Cons

- inefficient for **sorted** values
- must examine $n/2$ elements on average for a value that **is** in the array
- must examine **all n elements** for a value **not** in the array

Searching a Sorted List

- linear search works perfectly on an **unsorted** array of values
- for **unsorted** values, it is the **only** practical searching algorithm

Searching a Sorted List

- linear search works perfectly on an **unsorted** array of values
- for **unsorted** values, it is the **only** practical searching algorithm
- linear search also works perfectly on a **sorted** array of values
- however, is it **not** the only search algorithm for sorted values
- furthermore, it is not a even a **good** algorithm in this case

Enhanced Linear Search

- if the elements of the array are sorted, we can do much better
- we can stop when any of three conditions is true:
 1. we find the item
 2. we reach the end of the list
 3. we reach an element **bigger** than the element we're searching for

2	3	5	11	27	23	29
---	---	---	----	----	----	----

Enhanced Linear Search

- if the elements of the array are sorted, we can do much better
- we can stop when any of three conditions is true:
 1. we find the item
 2. we reach the end of the list
 3. we reach an element **bigger** than the element we're searching for

2	3	5	11	27	23	29
---	---	---	----	----	----	----

- search for the value 11
- do a **comparison** on 2, 3, 5, 11 (4 comparisons)
- return 3 (the position where 11 was found)

Enhanced Linear Search

- if the elements of the array are sorted, we can do much better
- we can stop when any of three conditions is true:
 1. we find the item
 2. we reach the end of the list
 3. we reach an element **bigger** than the element we're searching for

2	3	5	11	27	23	29
---	---	---	----	----	----	----

- search for the value 11
- do a **comparison** on 2, 3, 5, 11 (4 comparisons)
- return 3 (the position where 11 was found)

- search for the value 12
- do comparisons on 2, 3, 5, 11, 27 (just 5 comparisons)
- return a not-found indicator

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle
 - 1.3 the elements to the **right** of middle

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle
 - 1.3 the elements to the **right** of middle
2. if the range of elements is **empty**, return not-found sentinel

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle
 - 1.3 the elements to the **right** of middle
2. if the range of elements is **empty**, return not-found sentinel
3. else if the searched-for value is the **middle** element, you've found it and you're done

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle
 - 1.3 the elements to the **right** of middle
2. if the range of elements is **empty**, return not-found sentinel
3. else if the searched-for value is the **middle** element, you've found it and you're done
4. else if the searched-for value is **smaller** than the middle element, repeat step 1 on the **left half**

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

1. divide the range of elements to search into 3:
 - 1.1 the very **middle** element
 - 1.2 the elements to the **left** of middle
 - 1.3 the elements to the **right** of middle
2. if the range of elements is **empty**, return not-found sentinel
3. else if the searched-for value is the **middle** element, you've found it and you're done
4. else if the searched-for value is **smaller** than the middle element, repeat step 1 on the **left half**
5. else repeat step 1 on the **right half**

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11
- do a comparison on 11 (1 comparison)

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11
- do a comparison on 11 (1 comparison)
- return 3 (the position where 11 was found) — done

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11
- do a comparison on 11 (1 comparison)
- return 3 (the position where 11 was found) — done
- search for 7

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11
 - do a comparison on 11 (1 comparison)
 - return 3 (the position where 11 was found) — done
-
- search for 7
 - do a comparison on 11, 3, 5 (3 comparisons)

Binary Search

2	3	5	11	17	23	29
---	---	---	----	----	----	----

- search for 11
- do a comparison on 11 (1 comparison)
- return 3 (the position where 11 was found) — done

- search for 7
- do a comparison on 11, 3, 5 (3 comparisons)
- report not-found — done

Binary Search Implementation

- as with linear search, Gaddis' implementation won't compile on clang
- see `program_8_2_modified.cpp`

Binary Search Analysis

- do the same simulation with binary search as with linear search
- 10,000 times, search for a random number 0 – 1,000
- results:

Total hits: 971 Total misses: 9029

Minimum comparisons: 1 Maximum comparisons: 10

Average number of comparisons: 9.8768

- remember the linear results:

Total hits: 930 Total misses: 9070

Minimum comparisons: 1 Maximum comparisons: 1000

Average number of comparisons: 954.747

- on average linear search takes **a hundred times** as many steps as binary search!
- binary search **never** takes more than ten comparisons

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**
- if the value is not in the array, clearly it takes n steps to decrease the search space to size zero and determine not-found

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**
- if the value is not in the array, clearly it takes n steps to decrease the search space to size zero and determine not-found
- binary search also starts with a search space of size n

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**
- if the value is not in the array, clearly it takes n steps to decrease the search space to size zero and determine not-found
- binary search also starts with a search space of size n
- with each comparison, the search space decreases **by half**

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**
- if the value is not in the array, clearly it takes n steps to decrease the search space to size zero and determine not-found
- binary search also starts with a search space of size n
- with each comparison, the search space decreases **by half**
- how many steps does it take to repeatedly divide a number by 2 until you get to zero?

Binary Search Analysis

- let's dig a little deeper into exactly how long the two searches take
- linear search starts with a **search space** of size n
- with each comparison, the search space decreases **by one**
- if the value is not in the array, clearly it takes n steps to decrease the search space to size zero and determine not-found
- binary search also starts with a search space of size n
- with each comparison, the search space decreases **by half**
- how many steps does it take to repeatedly divide a number by 2 until you get to zero?
- the term for this is **logarithm**

Logarithms

- logarithms are particularly important in algorithm analysis
- computer scientists deal mainly with base-2 logarithms
- the simplest working definition of a base-2 logarithm is

Base-2 Logarithm

How many times can you divide a number n by 2 using integer division before the result is 1 or 0?

- the answer to this question is approximately the base-2 logarithm of n
- you can estimate base-2 logarithms directly from the powers of 2 table

Powers of 2

you should learn all the powers of 2 from 0 to 10

$$2^0 = 1$$

$$2^1 = 2$$

$$2^2 = 4$$

$$2^3 = 8$$

$$2^4 = 16$$

$$2^5 = 32$$

$$2^6 = 64$$

$$2^7 = 128$$

$$2^8 = 256$$

$$2^9 = 512$$

$$2^{10} = 1024$$

Binary Search

- remember the simulation results searching in an array with 1,000 elements
- the **most** comparisons ever needed was **10**
- the base-2 logarithm of 1,000 is approximately **10**
- the number of comparisons needed for binary search for an array of size n is at most $\log_2 n$

Binary Search

- remember the simulation results searching in an array with 1,000 elements
- the most comparisons ever needed was 10
- the base-2 logarithm of 1,000 is approximately 10
- the number of comparisons needed for binary search for an array of size n is at most $\log_2 n$
- compare this to linear search, which needs at most 1,000 comparisons for a 1,000-element array

Binary Search

Binary Search Pros

- much more efficient than linear search
- requires at most $\log_2 n$ comparisons
- easy algorithm to understand and code correctly

Binary Search Cons

- requires array elements to already be sorted

Binary Search

Binary Search Pros

- much more efficient than linear search
- requires at most $\log_2 n$ comparisons
- easy algorithm to understand and code correctly
- clearly, sorting is an issue, and we turn to that next

Binary Search Cons

- requires array elements to already be sorted