

Text Files and Strings

Class 24

Files

- a file on disk is strictly a sequence of bytes
- when you ask the operating system for stuff from a file, you just get raw bytes
- it is up to the programmer how to **interpret** those bytes

Files

- a file on disk is strictly a sequence of bytes
- when you ask the operating system for stuff from a file, you just get raw bytes
- it is up to the programmer how to **interpret** those bytes
- there are two main flavors of file, text and binary
- **every** file is a binary file in the sense that it contains bytes
- text files, however, contain **only** bytes that correspond to ASCII characters
- one of those bytes represents the **newline** character, interpreted as the end of a line
- thus text files are a sequence of **lines** each of which is a sequence of **characters**
- and we normally deal with text files **line-by-line**
- binary files must be done **byte-by-byte**


Values

- a digital computer can only manipulate bits, 0s and 1s
- a group of 8 bits is a byte
- a byte can be interpreted as an unsigned integer value in the range 0 – 255
- thus, the only fundamental values a computer can manipulate are the 256 values in the range 0 – 255
- at the lowest level, there are no characters, or doubles, or negative numbers, or strings
- there are only the binary values 0000 0000 – 1111 1111

Characters

- to **represent** a character, a program must use an **encoding**
- an encoding is an agreement about which bit pattern will represent which character
- e.g., let us agree that in a character context the byte 0100 0001 (which is 65_{10} or $0x41$) will represent 'A'
- C uses the ASCII encoding scheme

Text Files



what we see in an editor

A	l	i	c	e	\n	B	o	b	\n	C	a	r	o	l	\n
---	---	---	---	---	----	---	---	---	----	---	---	---	---	---	----

what is really in the file on disk

- this is why **every** line of output must be terminated with a newline

Streams

- the model that C uses for working with files is to consider them as **streams** of bytes
- to **read** from a file is to treat the file as an **input** stream of bytes coming in from disk
- to **write** to a file is to treat the file as a destination for an **output** stream of bytes going out to disk
- all files are accessed via a pointer to a FILE struct, which is defined in stdio.h

Filenames

- files on disk are identified by **filename**
- by convention a filename consists of **name** and **extension**
- e.g., `grades.xls` or `phone_plan.cpp`
- to access a disk file it must first be **opened**
- to open a file means to associate it with a variable which is a pointer to a `FILE` struct
- this is done using the `fopen` function

```
#include <stdio.h>
```

```
...
```

```
FILE* input_file = fopen("foo.bar", "r");
```

- can also open a file for writing ("`w`")
- this destroys contents if they exist

Reading and Writing

- it is possible to open a file for both reading and writing
 - "r+" read and write; don't overwrite contents first
 - "w+" read and write, and overwrite contents first
- this can be confusing
- much more potential for data corruption if something goes wrong
- not recommended; we won't do it

Closing a File

- before your program terminates
- you must **close** the file
- this frees up the operating system resources associated with the file
- for output files especially, this flushes the **write buffer** and ensures that everything you tried to write to the file is actually stored on disk

```
fclose(input_file);
```

Writing to Text Files

- writing to a text file is very safe
- `putc` — write a single character
- `fputs` — write a string and a newline
- `fprintf` — write a formatted string; often used by default, even if no formatting is to be done

Reading from Text Files

- reading input is fraught with danger
 - malicious data may have been introduced
 - even with no bad guys, data may not be in correct format
- `getc` — reads a single character, always safe, check for `feof`
- `fscanf` — can be useful, but is unsafe if data are not in the expected format
- `fgets` — the safe way to read a string because it has a size limit might corrupt data, but won't cause buffer overflow if the limit is correct

End of File

- typically we don't know how many lines of data a file contains
- typically we run `fgets` in a loop, stopping when it returns `NULL`
- sometimes you wish to check for end-of-file directly, with `feof()`
- `feof` only returns true **after** an attempted read has failed

Standard Streams

- every Unix program has three file streams automatically opened: `stdin`, `stdout`, and `stderr`
- `printf("foo")` is just a macro definition that really means `fprintf(stdout, "foo")`
- because the three standard files are handled automatically by the operating system, it is a semantic error to open or close them

fgets

- because fgets is the safest input routine, we should clearly understand how it works

```
#define MAXLINE 5  
char line[MAXLINE];  
fgets(line, MAXLINE, infile);
```

- fgets reads in at most **one less** than MAXLINE characters from the stream
- EOF and newline also cause reading to stop
- if a newline is read, it is stored into the string with the read pointer on the first character of the next line
- a terminating null character '\0' is **always** stored after the last character read into the string

Example

```
1 #include <stdio.h>
2 #define MAXLINE 5
3 int main(void)
4 {
5     char line[MAXLINE];
6     unsigned iter = 0;
7     while (fgets(line, MAXLINE, stdin))
8     {
9         printf("Iter %u: %s\n", iter++, line);
10    }
11    return 0;
12 }
```

123

1234

12345

123456

1234567

Definitions

```
#define MAXLINE 5  
char line[MAXLINE];
```

- how many “true” characters can line hold? only 4
- easy to get confused
- it is common to see definitions like this:

```
#define MAX_LINE_CHARS 5  
char line[MAX_LINE_CHARS + 1];
```

- but remember, if planning to use line with fgets, may need to account for newline **and** null character:

```
#define MAX_LINE_CHARS 5  
char line[MAX_LINE_CHARS + 2];
```