

Tables and Data

Class 29

Data Types

- databases are optimized to store data efficiently
- the type of data is very significant
- the broad categories are
 - character data
 - numeric data
 - time data
- MySQL does not have a Boolean data type
tinyint unsigned is used instead

Character Types

- `varchar(n)` stores up to n characters ($n \leq 255$)
 n is always 255; no point in any other value
- `char(n)` stores exactly n characters, right-padded with spaces
 to n characters ($n \leq 255$)
- these match HTML `<input type="text" />`
 - and other input types (e.g., phone, password, etc.)

Character Types

- text types
 - `tinytext` up to 2^8 bytes
 - `text` up to 2^{16} bytes
 - `mediumtext` up to 2^{24} bytes
 - `longtext` up to 2^{32} bytes
- these match HTML `<textarea>`

Character Types

enum a string-like object with a value chosen from a list of permitted values that are enumerated explicitly in the column specification at table creation time

- very compact storage, while simultaneously
- utilizing readable queries and results
- matches `<input type="text" />`

```
create table shirt
(
  name varchar(255),
  size enum('x-small', 'small', 'medium',
            'large', 'x-large')
);
```

Integer Types

`tinyint` 2^8 values

`smallint` 2^{16} values

`mediumint` 2^{24} values

`int` 2^{64} values

`bigint` 2^{128} values

- the default is signed
- use **unsigned** for unsigned things e.g., counts
- use **unsigned auto-increment** for unique keys
- these match HTML `<input type="number" />`

Exact Fixed-Point Types

`decimal(n,m)` stores n digits, m of which are to the right of the decimal point ($n \leq 65$)

`salary decimal(7, 2)`

can store every value from -99999.99 to 99999.99 with no approximation error

Floating Point Types

`float` uses 4 bytes for storage (rarely used)

`double` uses 8 bytes for storage

Time Types

date stores dates in yyyy-mm-dd format from 1000-01-01 to 9999-12-31

time stores times in hh:mm:ss format from -838:59:59 to 838:59:59

datetime stores times in yyyy-mm-dd hh:mm:ss format

- these match HTML `<input type="text" />`

Creating a Table

- create a simple table from the command line

```
create table pet
(
    id int unsigned not null auto_increment,
    name varchar(255),
    breed varchar(255),
    sex enum('F', 'M', 'U'),
    birth date
);
```

- but for reproducibility use a .sql file, e.g., create_pet.sql
- comments begin with “– ”

- use **source** command to read in the file

```
mysql> source create_pet.sql
```

- beware column names that are reserved words (e.g., order)

Entity Tables

entity tables store fundamental objects in your system
the equivalent of a **class**

- the name is a singular concrete noun
- each table row is one object
- columns store only singleton data (e.g., birthdate)
- no multi-value data (e.g., address) or derived data (e.g., age)
- each row is identified by a unique primary key (auto-increment or inherent)
- examples:
 - person
 - product
 - song
 - book

Primary Key

- each row of an entity table must have a **primary key**
- unique across all rows, not null, never changes
- some entities have natural primary keys (ISBN for a book, banner ID for a student)
- but **never** SSN for a person
- an entity without a natural primary key uses an auto_increment column

- example:

```
create table pet
(
  id int unsigned not null auto_increment,
  name varchar(255),
  breed varchar(255),
  sex enum('F', 'M', 'U'),
  birth date,
  primary key(id)
);
```

One-to-Many Tables (Attributes)

one-to-many tables store attributes of entities when an entity can have multiple values of an attribute

- the name is a concatenation of the entity name followed by the attribute name
- the attribute name is a concrete noun, abstract noun, sometimes derived from an intransitive verb
- **not** used when an attribute is itself an entity
- the entity is represented by its primary key
- pay attention to which columns can and cannot be null
- examples: `person_address`, `player_score`, `pet_vaccination`, `plant_growth`

One-to-Many Attribute Tables

```
create table pet_vaccination
(
    id int unsigned not null auto_increment,
    pet_id int unsigned not null,
    name varchar(255) not null,
    vet varchar(255),
    administered date,
    primary key(id),
    constraint pet_fk foreign key(pet_id)
        references pet(id)
);
```

Redundant Data

- a key concept of RDBMS is that every datum is stored in exactly one place
- there are no duplicate or redundant data in the system
- so, if we store birthdate, we do not store age, because that is redundant

Duplicate Data

- consider the words.txt file:

abate	verb	to put an end to; nullify
abeyance	noun	a lapse in succession
abjure	verb	to renounce upon oath
abrogate	verb	to abolish by authoritative
abstruse	adjective	difficult to comprehend
acarpous	adjective	effete; no longer fertile

- the parts of speech are duplicate data
- even without the parts.txt file, but especially if that file exists

Eliminating Duplicate Data

instead we need three tables

id	part	id	word	definition
1	adjective	1	abate	to put an end to; nullify
2	adverb	2	abeyance	a lapse in succession
3	noun	3	abjure	to renounce upon oath
4	verb	4	abrogate	to abolish by authoritative
		5	abstruse	difficult to comprehend
		6	acarpous	effete; no longer fertile

word_id	part_id
1	4
2	3
3	4
4	4
5	1
6	1

Entity-Entity Tables

- express a relationship between two entities
- may be one-to-one (a person has one biological mother)
- may be one-to-many (a person places many orders)
- may be many-to-many (a physician has many patients, and a patient may see many physicians)
- the name is the concatenation of the two entity names (e.g., person_order)
- or the name of the relationship (e.g., role for the relationship between actor and movie)
- columns consist of the primary keys of the two entities, and possibly other columns which are attributes of the relationship

Entity-Entity Tables

example many-to-many table

```
create table physician_patient
(
    physician_id int unsigned not null,
    patient_id int unsigned not null,
    first_visit date default null,
    primary key(physician_id, patient_id),
    constraint physician_fk foreign key(physician_id)
        references person(id),
    constraint patient_fk foreign key(patient_id)
        references person(id)
);
```

- no auto_increment column
- note the foreign key constraints

Entity-Entity Tables

Example one-to-many table

```
create table movie_director
(
    director_id int unsigned not null,
    movie_id int unsigned not null,
    salary decimal(8,0) default null,
    primary key(director_id, movie_id),
    unique key(movie_id),
    constraint director_fk foreign key(director_id)
        references director(id),
    constraint movie_fk foreign key(movie_id)
        references movie(id)
);
```

Inserting Data

data are inserted into only one table at a time

```
insert into actor  
  (first_name, last_name, sex, film_count)  
values  
  ('Shailene', 'Woodley', 'F', 13);
```

- any attributes (columns) left out explicitly get set either to the default value (if any) or to NULL
- an auto-increment column automatically gets set to the next-highest value for that column
- it is a logical error to leave out a column that is declared as not null and which has no default value

Retrieving Information: One Table

a query that involves only one table is easy:

- what columns have the information you need?
- do any column data need to be transformed?
- what rows have the information you need?
- do any rows need to be aggregated?
- in what order should the information be delivered?

Example: can we see the movies of the 90's listed by rank?

```
select name, rank
from movie
where year between 1990 and 1999
order by rank desc;
```

Two Tables

- some queries require information from two different tables
- this is where a DB blows away text files

Example: the names of comedy movies

- the **movie** table has names
- the **movie_genre** table has genres
- the two tables need to be **joined** to get both pieces of information at the same time
- almost always, the join condition will involve primary and foreign keys

```
select movie.name
from movie
join movie_genre on movie.id = movie_genre.movie_id
where movie_genre.genre = 'Comedy';
```

Ambiguous Column Names

- as soon as more than one table is involved, column names can be ambiguous
- for example, movie, actor, and director each have a column named "id"
actor and director both have columns named "first_name" and "last_name"
- even if right now there's only one table with a column named "rank" (movie), tomorrow someone might add a new column in, say, the actor table named "rank" and then "rank" would be ambiguous
- therefore, every column name in a query with more than one table must be fully qualified

```
select movie.name  
from movie  
join movie_genre on movie.id = movie_genre.movie_id  
where movie_genre.genre = 'Comedy';
```


Shortcut Table Names

- fully qualified table names can become tedious
- so within each query, it is common to use shortcut names

```
select m.name  
from movie as m  
join movie_genre as mg on m.id = mg.movie_id  
where mg.genre = 'Comedy';
```

Two Tables

the names of directors who have a better than even chance of directing a comedy

- the **director** table has directors' names
- the **director_genre** table has genre and probability information

```
select concat_ws(' ', d.first_name, d.last_name) as name,  
       format(dg.prob, 3) as prob  
from director as d  
join director_genre as dg on d.id = dg.director_id  
where dg.genre = 'Comedy' and dg.prob > 0.5  
order by dg.prob desc;
```

Three Tables

the names of actresses and their characters who played in movies in the 1990s

- the names of actresses are in the actor table
- the names of characters are in the role table
- the year of movies is in the movie table

```
select concat_ws(' ', a.first_name, a.last_name) as actorname,  
       r.role  
from actor as a  
join role as r on a.id = r.actor_id  
join movie as m on r.movie_id = m.id  
where m.year between 1990 and 1999  
       and a.sex = 'F'  
order by a.last_name asc;
```

Four Tables

the names of actresses and their characters who played in comedies in the 1990s, along with the name of the movie

- the names of actresses are in the actor table
- the names of characters are in the role table
- the year and title of movies are in the movie table
- the genre of movies is in the movie_genre table

```
select concat_ws(' ', a.first_name, a.last_name) as actname,  
       r.role, m.name  
from actor as a  
join role as r on a.id = r.actor_id  
join movie as m on r.movie_id = m.id  
join movie_genre as mg on mg.movie_id = m.id  
where m.year between 1995 and 2005  
      and a.sex = 'F'  
      and mg.genre = 'Comedy'  
order by a.last_name asc;
```