

C.R.U.D. in MongoDB

**Dr. Chetan Jaiswal,
Department of Computer Science,
Truman State University
cjaiswal@truman.edu**

C.R.U.D.

- Create – INSERT (Mongo) - INSERT (SQL)
- Read – FIND (Mongo) – SELECT (SQL)
- Update – UPDATE (Mongo) – UPDATE (SQL)

One variation in Mongo is UPSERT

- Delete – REMOVE (Mongo) – DELETE (SQL)

C.R.U.D.

- Mongo Shell is simple javascript shell:

```
for(i=0;i<5;i++) print("Hello" + i);
```

- Type “help” for help tutorials
- Using built-in functions: BSON (<http://bsonspec.org/>)

```
obj={a:1, b:2, c:"Hello", d:["Apples", "Oranges"],e:new Date()}
```

INSERT

- Let document be:
 - Use “insert” command
 - Use mongoimport

```
{ author: 'joe',  
  title : 'Yet another blog post',  
  text : 'Here is the text...',  
  tags : [ 'example', 'joe' ],  
  comments : [  
    { author: 'jim',  
      comment: 'I disagree'  
    },  
    { author: 'nancy',  
      comment: 'Good post'  
    }  
  ]  
}
```

- Do it now:

Insert a document into the *fruit* collection with the attributes of "name" being "apple", "color" being "red", and "shape" being "round". Use the "insert" method.

findOne

- It is used to get one document out of the collection
- `db.test.findOne(<Conditions>, <Projections>)`
- `db.test.findOne({name:"James"}, <Projections>)`
- `db.test.findOne({name:"James"}, {name:true, _id:false})` // by default `_id` will always be present until explicitly specified "false"
- `db.zips.distinct("state")`

find

- It is used to get documents out of the collection
- `db.test.find(<Conditions>, <Projections>)`
- `db.test.find({name:"James"}, <Projections>)`
- `db.test.find({name:"James"}, {name:true, _id:false})` // by default `_id` will always be present until explicitly specified "false"
- `db.test.find({name:"James", zip:64110}, {name:true, _id:false})`
- Assume documents to be like:
- how would you find all documents with type: essay and score: 50 and only retrieve the *student* field?

```
{
  "_id" : ObjectId("50906d7fa3c412bb040eb577"),
  "student_id" : 0,
  "type" : "exam",
  "score" : 54.6535436362647
}
```

find with \$gt/\$gte/\$lt/\$lte

- It is used to get documents based on range of values
- `db.test.find(<Conditions>, <Projections>)`
- `db.test.find({score:{$gt:50}}, <Projections>)`
- `db.test.find(({score:{$gt:50, $lte:100}}, {name:true, _id:false}))` // by default `_id` will always be present until explicitly specified “false”
- Can also be used with strings
- `db.test.find({name:{$gt:"D"}})` //lexicographically using UTF8
- `db.test.find({name:{$gt:"D", $lte:"E"}})` // will this document be part of the result?:
`{_id:87, name: "DOG", zip: 64110}` Unicode Transformation Format - 8 unicode characters
- Can name field have a number instead of string? Like: `{_id:87, name: 60, zip: 64110}`
- How will this query work now?: `db.test.find({name:{$gt:"D", $lte:"E"}})`
- The MongoDB comparisons do not span between data types

regex, \$type, \$exists

- \$exists: allows you to query documents which contains a particular attribute
`db.people.find({profession: {$exists:true}})// or false`
- \$type: allows you to query documents where a particular field is of certain data type.

`db.people.find({name:{$type:2}})//2 corresponds to string type`

Refer: <https://docs.mongodb.org/v3.0/reference/operator/query/type/>

`>db.posts.find({post_text:{$regex:"icecream"}})`

`>db.posts.find({post_text:/icecream/})`

`>db.posts.find({post_text:{$regex:"icecream",$options:"$i"}})//case insensitive`

`>db.products.find({ sku: { $regex: /789$/ } })//ends with 789`

`>db.products.find({ sku: { $regex: /^ABC/i } })//starts with ABC`

\$or/\$and

- `db.people.find({ $or: [{name: { $regex: "e$" } }, { age: { $exists:true } }] })`
- \$and can be used similarly
- How would you find all documents in the `scores` collection where the `score` is less than 50 or greater than 90?
- How about this: `db.scores.find({ score : { $gt : 50 }, score : { $lt : 60 } })`

Querying inside arrays

- Same as normal documents
- Can use \$in and \$all
- `db.accounts.find ({ favorites: { $all: ["pretzels", "beer"] } })` // finds all the docs where the array favorites contain pretzels and beer
- `db.accounts.find ({ name: { $in: ["John", "Howard"] } })` // find all the docs where name contains John or Howard
- `db.accounts.find ({ favorites: { $in: ["beer", "icecream"] } })`???

\$set/\$unset

- Assume a doc like: { _id:..., name: "Helen" }
- We would like to give age to Helen. How can we do that?
- a. db.people.update({ name: "Helen" }, { name: "Helen", age: 17.5 })

Is this way good? Wouldn't it delete everything else from the previous doc?

- b. db.people.update ({ name: "Helen" }, { \$set: { age: 17.5 } })

This will just add one more field in the doc without disturbing the previous fields.

- \$unset is just the opposite

More \$ operators

- Assume a doc: { _id: 0, a: [1, 2, 3, 4] }
- `db.collection.update ({ _id:0 }, { $set: { "a.2": 5 } })`
- `db.collection.update ({ _id:0 }, { $push: { a: 6 } })`
- `db.collection.update ({ _id:0 }, { $pop: { a: 1 } })` //right most
- `db.collection.update ({ _id:0 }, { $pop: { a: -1 } })` //left most
- `db.collection.update ({ _id:0 }, { $pushAll: { a: [5, 6, 7] } })`
- `db.collection.update ({ _id:0 }, { $pull: { a: 5 } })` //removes the value 5
- Similarly \$pullAll
- Array can be treated as SET
- `db.collection.update ({ _id:0 }, { $addToSet: { a: 8 } })` //a will be now treated as a SET, its idempotent, meaning if you run above query multiple times it will only add 8 once

<https://docs.mongodb.com/manual/reference/operator/update/set/>

remove

- remove: it deletes the document specified
- `db.people.remove({ name: "Alice" })`
- `db.people.remove({ name: { $gt:"M" } })`
- How about: `db.people.remove()???`
- Instead give: `db.people.remove({ })`
- Or you can do: `db.people.drop()`
- Which is better?

Mmapv1 engine

- Mmap is a system call (refer mmap MAN page)
- File is mapped to exact sized virtual memory
- Provides collection level locking
- In place update

WiredTiger engine

- It is faster for a lot of workloads
- Document level concurrency/ no locking engine
- If two writes come for same document, one is rolled back and restarted again later
- Offers compression and indexes
- No inplace update
- Default engine is mmapv1: `mongod --dbpath <some-new-directory> -storageEngine wiredTiger`

indexes

- Faster reads and slower writes
- You can create composite/compound index
- `db.test.ensureIndex ({ a: 1 })` // 1 represents ascending
- Same as using `db.test.createIndex`
- `db.students.createIndex({"class":1,"student_name":1})`
- `db.collection.getIndexes()`
- `db.collection.dropIndex({ a:1 })`
- `db.students.createIndex({student_id:1,class_id:1},{unique:true})`
- `db.students.createIndex({student_id:1,class_id:1},{unique:true, sparse:true})`//some docs may not have the field on which you are trying to create index on

Using explain

- `db.collection.explain().find()` // explains how the query will be processed
- `db.collection.explain().help()`
- `db.example.example.find({ a:5, b:23 })`, to know how mongodb is processing the query run with `explain`.
- `db.example.explain().find({ a:5, b:23 })`
- `db.example.createIndex({ a:1 })` // run `explain` after this and see if it is getting used
- `Explain` can be used in 3 different modes: `queryPlanner`, `executionStats`, `allPlansExecution`
- Latter ones include the former ones!
- `db.example.explain("executionStats").find({ a:5, b:23 })`
- `allPlansExecution`: this mode shows what query optimizer has done, that is to explore all possible options
- Covered Queries: MongoDB programmers strive to satisfy all the front end queries using indexes. This makes the queries a lot faster. Use `explain` to find if the query is covered or not.

```
for(i=0;i<500000;i++) db.test.insert({a:i,b:500000-i,c:i%46})
```

Aggregation – Self Study

- Refer: <https://docs.mongodb.org/manual/core/aggregation-introduction/>